

THOMAS JIRALERSPONG

☎ +1 514 625 9308 | @ thomas@jiralerspong.com |  LinkedIn |  GitHub |  Website |  Google Scholar |  Games

EDUCATION

Université de Montréal/Mila

PhD – Computer Science – Supervised by [Yoshua Bengio](#), [Guillaume Lajoie](#) & [Doina Precup](#)

2023 – Present

EXPERIENCE

Astra Fellowship/MATS - Mentored by [Dan Mossing](#) (Anthropic)

Research Fellow

Apr 2026 – Aug 2026

- **Project:** Relationship between LLM context and answer representations.

Astra Fellowship - Google DeepMind Stream

Research Fellow

Jan 2026 – Mar 2026

- **Project:** [Control intervention awareness benchmark](#), receiving the ICLR 2026 CAO workshop Best Paper Award.

Anthropic - Mentored by [Trenton Bricken](#)

Research Fellow

May 2025 – Oct 2025

- **Project:** Cross-architecture crosscoder model diffing: [official Anthropic research release](#) (500K+ views).

Occam AI

Research Scientist - New York City, United States

Jun 2024 – Apr 2025

- **Project:** Optimization methods for multi-agent LLM systems and automated SQL-generation pipelines.

AWARDS, GRANTS & SCHOLARSHIPS

Manifund research grants (US\$473,332) for mentored AI safety and interpretability research projects: [\[1\]](#), [\[2\]](#).

Vanier Scholarship (150 000\$), FRQNT Scholarship (40 000\$) – Rank #1, NSERC Scholarship (17 500\$)

[Y Combinator \(W26\)](#) – Selected from 20,000+ applicants (<1% acceptance rate); declined to focus on PhD. 2025

Chosen as one of the 200 most promising young researchers in math & CS by the [Heidelberg Laureate Forum](#) 2023

Winner of [Project X 2021](#) (ML Research Competition - 25 000\$) 2021

RESEARCH LEADERSHIP & MENTORSHIP

SPAR – Research Mentor, 3 AI safety projects

Aug 2026 – Present

AlgoVerse AI Research – Principal Investigator, oversight of 15 projects

Jun 2026 – Sep 2026

MARS V - Meridian Cambridge – Research Mentor, 3 projects

Jun 2026 – Present

Mila/McGill University – Research Mentor, 4 researchers (2 PhD, 1 MSc, 1 intern)

2023 – Present

FIRST-AUTHORED PUBLICATIONS

[Measuring Control Intervention Awareness Across Frontier LLMs](#)

J. Schaeffer*, [T. Jiralerspong](#)*, A. Panfilov, R. Zimmermann

Best Paper Award, Oral Presentation

CAO Workshop at ICLR 2026

[Noticing the Watcher: LLM Agents Can Infer CoT Monitoring from Blocking Feedback](#)

[T. Jiralerspong](#)*, F. Kondrup, Y. Bengio

AIWILD Workshop at ICLR 2026

[Cross-Architecture Model Diffing With Crosscoders: Unsupervised Discovery of Differences Between LLMs](#)

[T. Jiralerspong](#), T. Bricken

Official Anthropic Release

[Learning What Matters: Steering Diffusion via Spectrally Anisotropic Forward Noise](#)

L. Scimeca*, [T. Jiralerspong](#)*, B. Earnshaw, J. Hartford, Y. Bengio.

Preprint, 2025

[A Complexity-based Theory of Compositionality](#)

E. Elmoznino*, [T. Jiralerspong](#)*, Y. Bengio, G. Lajoie.

ICML 2025

[Geometric Signatures of Compositionality Across a Language Model's Lifetime](#)

J. Lee*, [T. Jiralerspong](#)*, L. Yu, Y. Bengio, E. Cheng.

Oral Presentation, Top 1% of papers

ACL 2025

*Equal Contribution

MENTORED PUBLICATIONS

[When Does Interleaving Prevent Emergent Misalignment?](#)

S. Mammeri, R. N. Kwon, C. Sun, S. Sahoo, C. Gagné, [T. Jiralerspong](#)

Mech Interp Workshop at ICML 2026

[Hidden-State Similarity Predicts Re-Elicitation After Inoculation Prompting](#)

R. N. Kwon, S. Mammeri, S. Sahoo, C. Gagné, [T. Jiralerspong](#)

Mech Interp Workshop at ICML 2026