

THOMAS JIRALERSPONG

[+1 514 625 9308](tel:+15146259308) | @ thomasjiralerspong@gmail.com | [LinkedIn](#) | [GitHub](#) | [Website](#) | [Google Scholar](#) | [Games](#)

Citizenships: Canadian and Italian (European Union)

EDUCATION

Université de Montréal/Mila

PhD – Computer Science – Supervised by Yoshua Bengio, Guillaume Lajoie & Doina Precup 2023 – 2027 (Expected)

- **Awards:** Vanier Scholarship (150 000\$), FRQNT Scholarship (40 000\$) (Rank #1), NSERC Scholarship (17 500\$)

Massachusetts Institute of Technology (MIT)

Brains, Minds & Machines Summer Course

2024

McGill University

B.Sc. – Honours Computer Science – GPA:4.00/4.00 – Supervised by Doina Precup & Blake Richards 2020 – 2023

EXPERIENCE

Astra Fellowship - Google DeepMind Stream

Research Fellow - Berkeley, United States

PyTorch, Python

Jan 2026 – Present

Anthropic

Research Fellow - San Francisco, United States

PyTorch, Python

Apr 2025 – Oct 2025

- **Project:** Mechanistic interpretability to discover behavioral differences between models

Occam AI

Research Scientist - New York City, United States

PyTorch, Python

Jun 2024 – Apr 2025

- **Projects:** Optimization of interactions between networks of LLM agents, SQL query generation using LLMs

Waabi

Deep Learning Research Intern - Toronto, Canada

PyTorch, Python

Jun 2023 – Sep 2023

- **Project:** Realistic and controllable traffic simulation using a transformer based variational autoencoder

Amazon Web Services (AWS) – S3 Team

Software Development Engineer Intern – Vancouver, Canada

Python, JavaScript

May 2022 – Jul 2022

Expedia Group

Software Development Engineer Intern – Montreal, Canada

JavaScript, TypeScript, React

Jun 2019 – Aug 2019

SELECTED PUBLICATIONS

Cross-Architecture Model Differing With Crosscoders: Unsupervised Discovery of Differences Between LLMs

T. Jiralerpong, T. Bricken

ICLR 2026 (Submitted)

Learning What Matters: Steering Diffusion via Spectrally Anisotropic Forward Noise

L. Scimeca*, T. Jiralerpong*, B. Earnshaw, J. Hartford, Y. Bengio.

ICLR 2026 (Submitted)

A Complexity-based Theory of Compositionality

E. Elmoznino*, T. Jiralerpong*, Y. Bengio, G. Lajoie.

ICML 2025

Geometric Signatures of Compositionality Across a Language Model's Lifetime

J. Lee*, T. Jiralerpong*, L. Yu, Y. Bengio, E. Cheng.

Oral Presentation, Top 1% of papers

ACL 2025

Efficient Causal Graph Discovery Using Large Language Models

T. Jiralerpong*, X. Chen*, Y. More, V. Shah, Y. Bengio

ICLR Workshop 2024

Towards Safe Mechanical Ventilation Treatment Using Deep Offline Reinforcement Learning

F. Kondrup*, T. Jiralerpong*, E. Lau*, N. de Lara, J. Shkrob, M.D. Tran, D. Precup, S. Basu.

AAAI 2023

*Equal Contribution

AWARDS & ACHIEVEMENTS

Y Combinator (W26) – Selected from 20,000+ applicants (<1% acceptance rate); declined to focus on PhD.

2025

Chosen as one of the 200 most promising young researchers in math & CS by the **Heidelberg Laureate Forum**

2023

Winner of **Project X 2021** (ML Research Competition - 25 000\$)

2021