

# THOMAS JIRALERSPONG

☎ +1 514 625 9308 | @ [thomasjiralerspong.com](https://twitter.com/thomasjiralerspong) | [LinkedIn](#) | [GitHub](#) | [Website](#) | [Google Scholar](#) | [Games](#)

**Citizenships:** Canadian and Italian (European Union)

## EDUCATION

### Université de Montréal/Mila

PhD – Computer Science – Supervised by [Yoshua Bengio](#), [Guillaume Lajoie](#) & [Doina Precup](#) 2023 – Present

- **Awards:** Vanier Scholarship (150 000\$), FRQNT Scholarship (40 000\$) (Rank #1), NSERC Scholarship (17 500\$)

## EXPERIENCE

### Astra Fellowship - Mentored by [Dan Mossing](#) (Anthropic)

Research Fellow

PyTorch, Python

Apr 2026 – Aug 2026

- **Project:** Characterizing the mapping from context to answer activations in LLMs

### Astra Fellowship - Google DeepMind Stream

Research Fellow

PyTorch, Python

Jan 2026 – Mar 2026

- **Project:** Benchmark for control intervention awareness in LLMs

### Anthropic - Mentored by [Trenton Bricken](#)

Research Fellow

PyTorch, Python

May 2025 – Oct 2025

- **Project:** Mechanistic interpretability to discover behavioral differences between models
- Published as [official Anthropic research release](#) (500K+ views)

### Occam AI

Research Scientist - New York City, United States

PyTorch, Python

Jun 2024 – Apr 2025

- **Projects:** Optimization of interactions between networks of LLM agents, SQL query generation using LLMs

## MENTORSHIP

**SPAR (Supervised Program for Alignment Research)** – Research Mentor, 3 AI safety projects

Aug 2026 – Present

**AlgoVerse AI Research** – Principal Investigator, overseeing 20 AI safety research teams

Jun 2026 – Present

**MARS V - Meridian Cambridge** – Research Mentor, 7 AI safety researchers (incl. PhD/MSc)

Jun 2026 – Present

**Mila/McGill University** – Research Mentor, 2 PhD students, 1 research intern & 1 Master's student

2023 – Present

## SELECTED PUBLICATIONS

### [When Does Interleaving Prevent Emergent Misalignment?](#)

S. Mammeri, R. N. Kwon, C. Sun, S. Sahoo, C. Gagné, [T. Jiralerspong](#)

*Mech Interp Workshop at ICML 2026*

### [Hidden-State Similarity Predicts Re-Elicitation After Inoculation Prompting](#)

R. N. Kwon, S. Mammeri, S. Sahoo, C. Gagné, [T. Jiralerspong](#)

*Mech Interp Workshop at ICML 2026*

### [Measuring Control Intervention Awareness Across Frontier LLMs](#)

J. Schaeffer\* [T. Jiralerspong\\*](#), A. Panfilov, R. Zimmermann

*Best Paper Award, Oral Presentation*

*CAO Workshop at ICLR 2026*

### [Noticing the Watcher: LLM Agents Can Infer CoT Monitoring from Blocking Feedback](#)

[T. Jiralerspong\\*](#), F. Kondrup, Y. Bengio

*AIWILD Workshop at ICLR 2026*

### [Cross-Architecture Model Diffing With Crosscoders: Unsupervised Discovery of Differences Between LLMs](#)

[T. Jiralerspong](#), T. Bricken

*Official Anthropic Release*

### [Learning What Matters: Steering Diffusion via Spectrally Anisotropic Forward Noise](#)

L. Scimeca\*, [T. Jiralerspong\\*](#), B. Earnshaw, J. Hartford, Y. Bengio.

*ICML 2026 (Submitted)*

### [A Complexity-based Theory of Compositionality](#)

E. Elmoznino\*, [T. Jiralerspong\\*](#), Y. Bengio, G. Lajoie.

*ICML 2025*

### [Geometric Signatures of Compositionality Across a Language Model's Lifetime](#)

J. Lee\*, [T. Jiralerspong\\*](#), L. Yu, Y. Bengio, E. Cheng.

*Oral Presentation, Top 1% of papers*

*ACL 2025*

\*Equal Contribution

## AWARDS & ACHIEVEMENTS

**Manifund research grants** – US\$473,332 for [mentored AI safety](#) and [alignment pretraining](#) research.

2026

**Y Combinator (W26)** – Selected from 20,000+ applicants (<1% acceptance rate); declined to focus on PhD.

2025

Chosen as one of the 200 most promising young researchers in math & CS by the [Heidelberg Laureate Forum](#)

2023

Winner of **Project X 2021** (ML Research Competition - 25 000\$)

2021